FAQ

1. *Should I submit the .ipynb file or the .py?*
    - Submit only the .ipynb file

2. *Should the size of the training and validation sets be equal, or should they be ⅔ - ⅓ ?*
    - They were meant to be equal (50-50), however it is fine to go for ⅔ - ⅓ if the point raised ambiguity
    - The important thing is to respect class balancing

3. *Should points 4 and 5 be implemented with transfer learning?*
    - Yes

4. *Should a criterion be used for searching hyperparameters?*
    - It is not required but it is preferable to adopt some kind of criterion and motivate it

5. *Should multiple code notebooks be provided (e.g. one for each point)?*
    - Provide just the last implementation, leave previous points as comments

FAQ

6. *Which hyperparameters can we optimize?*
   - Learning rate, number of epochs and step down policy are good candidates
   - Batch size and weight decay are also options

7. *Should results be fully deterministic?*
   - No. If you observe significant deviation between experiments, do multiple runs and report average and std

8. *I have used hyperparameters X/Y and the result is Z, is this ok?*
   - Annotate the results you obtained and write them in the final report

9. *My answer for point A is this, is it ok?*
   - Provide your implementation in the code, you will get feedback after the hard deadline

FAQ

10. *Technical issues (colab notebook, python errors)*
     - Many problems with colab can be solved by restarting the runtime (Runtime -> Restart and run all)
     - Python bugs, pytorch errors, colab issues and other problems are considered as part of the assignment